

Supplementary Materials

Matthew Brown and Harsh Gupta

July 27, 2026

A Materials and Methods

A.1 Media outlet slant

We measure outlet slant using the estimates from [Bakshy et al. \[2015\]](#); henceforth BMA. The BMA score reflects the self-reported ideology of an outlet’s online audience. Each article produced by an outlet is given a score from -2 to 2, depending on the average ideology of users sharing it on Facebook. Articles shared mainly by liberals get negative scores, while those shared by conservatives receive positive scores. The outlet’s BMA score is the average of all its articles’ scores.

A.2 Citations of research by media outlets

We use Altmetric to construct our media citation data. Altmetric is a bibliometric database that scrapes the websites of thousands of major news outlets every day, counting every instance that a scientific paper is mentioned. Altmetric began

scraping news sites in early 2013, and we conducted our latest search in 2025.

Most citations identified by the Altmetric scraping tool are from news articles that include hyperlinks to scientific journal websites papers or DOIs. A minority of citations are also identified via text mining. For a citation to be identified via text mining, the news article must include the scientist’s name, a publication date, and the name of the publishing journal. Some prominent news outlets have few observations or are missing from the Altmetric database (for example, the Wall Street Journal). The reason is that Altmetric cannot scrape websites with restricted content such as paywalls.¹

Many news organizations publish stories written by journalists at independent organizations, such as the Associated Press (AP). An article from an AP reporter about a topic of national interest might appear in dozens of local newspapers. Since individual news outlets have relatively little agency over the content that appears in these syndicated articles, we restrict our sample to articles that are only published in a single news outlet. In particular, we drop all observations for which multiple news articles share an identical title.

We *a priori* exclude two websites—*globalresearch.ca* and *iflscience.com*—that specialize in science coverage. We also drop news outlets that cover fewer than 20 scientific articles. Our final sample consists of 139 media outlets.

¹Source: Personal correspondence with Altmetric

A.3 Scientific papers

Our primary data source for scientific articles is Dimensions [[Herzog et al., 2020](#)]. Dimensions is a large bibliometric database that aggregates information on published scientific research across the world. Using their API, we downloaded the scientific papers published each day from 2015 to 2023.²

For each paper, we collect data on authors: names, university affiliations, fields, and locations. We also collect paper-level variables such as field of research, journal of publication, and year of publication. Last, we collect information on different quality measures of research: number of scientific citations, journal impact, research funding, and patent citations.

A.3.1 Construction and Interpretation of Quality Index

In our analysis, we use the aforementioned measures to capture the scientific quality of an article. In general, it is challenging to define and measure the quality of science. The key reason is that the quality of a scientific article is multifaceted. For example, scientific quality can be defined based on impact, plausibility, novelty, practical applicability, or relevance. Different measures captures different notions of quality. Any single measure of quality will likely miss some important notions of quality.

Thus, we combine these different measures into an index using a principal component analysis (PCA). Our index is based on the number of scientific citations, the

²Due to rate limitations, we restricted our query to 1000 articles each day. In the rare instances with more than a thousand articles published in a day, we include the 1000 scientific articles with the most academic citations.

quality of the journal, the research funding, and the patent citations for a scientific article. These measures capture quality from different perspectives. Scientific citations capture relevance from the perspective of other scientists. Research funding captures potential value from the perspective of the sponsor. Patent citations capture the practical applicability of a scientific article. Our index is the first principal component of this array of measures.

We contend that these measures capture a wide range of quality criteria and our index provides a low-dimensional variable to capture these criteria. We emphasize that we are constructing a statistical measure (and not a normative one) that helps us coarsely distinguish between low and high quality research. We do not take a normative stance on this measure, but rather aggregate commonly used and widely accepted measures of scientific quality.

A.4 Political Ideology of Individuals

We use the Database on Money, Ideology, and Elections (DIME) [Bonica, 2024] to construct ideology scores for scientists. DIME compiles public data on political contributions to U.S. politicians at the federal, state, and local levels from 1976–2014. The donation behavior of each contributor is summarized by a *Campaign Finance (CF) Score*, which places the contributor on a one-dimensional left-right scale Bonica [2014].

A.5 Linking Scientists to Campaign-Finance Records

We now describe the procedure we use to link the scientists in our Dimensions sample to individual donors in the Database on Ideology, Money in Politics, and Elections [DIME; Bonica, 2014], from which we obtain each scientist’s Campaign Finance (CF) score. Our approach follows the canonical probabilistic record-linkage model of Fellegi and Sunter [1969] as implemented at scale by Enamorado et al. [2019], and adopts best practices from Abramitzky et al. [2021].

A.5.1 Data sources and target populations

The linkage joins two files. The *scientist* file is the set of authors who published at least one article indexed by Dimensions between 2015 and 2024 with a U.S. affiliation that was covered by at least one media outlet. For each author we retain every reported affiliation over the sample period, so that an author is represented by the *set* of states and institutions at which she published. The *donor* file is constructed from DIME. We restrict to individual contributors (`contributor.type = I`) whose most recent contribution falls in or after the 2006 cycle and whose self-reported occupation indicates an academic or research vocation (“professor,” “scientist,” “researcher,” “student,” “physician,” or “doctor”). Crucially, we draw on the *transaction-level* DIME contribution records rather than the contributor-level summary, which allows us to represent each donor by the set of states, employers, and cycles observed across all of her contributions.

A.5.2 Blocking

Comparing all possible matches is infeasible. We restrict comparisons to candidate pairs that agree on a phonetic match of the last name, which is robust to spelling and transcription variants. We run blocking based on both the NYSIIS and Double Metaphone encodings so that pairs missed by one key are recovered by the other.

A.5.3 Probabilistic linkage model

For each candidate pair of records (i, j) we model the probability that the pair is a true match. The match probability depends on the alignment of certain fields (characteristics) between the pair. Specifically, we compare first and last names by the Jaro–Winkler string similarity, discretized into three levels (disagreement, partial agreement, and agreement); the middle initial, gender, and a geographic indicator. Let γ_{ij} denote the resulting agreement pattern and $M_{ij} \in \{0, 1\}$ the latent match status. (For example, $\gamma_{ij}^{M1} = 1$ if the middle initials are the same.) Under the conditional-independence assumption of the [Fellegi and Sunter \[1969\]](#) model, the data are a mixture

$$\Pr(\gamma_{ij}) = \lambda \prod_f \Pr(\gamma_{ij}^f \mid M_{ij} = 1) + (1 - \lambda) \prod_f \Pr(\gamma_{ij}^f \mid M_{ij} = 0),$$

where f indexes fields, λ is the marginal match rate, and the field-level probabilities $m_f = \Pr(\gamma^f \mid M = 1)$ and $u_f = \Pr(\gamma^f \mid M = 0)$ are estimated by the Expectation–Maximization algorithm [[Enamorado et al., 2019](#)].

We apply a term-frequency adjustment to the name fields, so that agreement on a rare name is weighted more heavily than agreement on a common name.

Academic careers are geographically mobile, and the two data sources record location at different moments: the donor’s state and employer are observed at the time of donation, whereas the scientist’s affiliation is observed at the time of publication. A true match may therefore disagree on location. Following the logic of [Abramitzky et al. \[2021\]](#), who link on time-invariant characteristics precisely because residence changes, we do not impose location agreement as a requirement. Instead, we exploit the fact that both records carry a *history*. For each candidate pair we compute whether the donor’s set of observed states intersects the scientist’s set of observed states, yielding a single categorical indicator (overlap, no overlap, or missing). Agreement constitutes positive evidence; disagreement carries only a mild penalty, as estimated by the EM procedure. Similarly, overlaps in employment history, rather than exact matching, is used as another characteristic.

The model returns, for each candidate pair, a posterior match probability $\zeta_{ij} = \Pr(M_{ij} = 1 \mid \gamma_{ij})$.

A.5.4 Match Threshold

The posterior ζ_{ij} partitions candidate pairs into accepted, rejected, and ambiguous regions. We select the acceptance threshold on ζ to control the false-discovery rate among accepted matches, using a manually labeled audit (Section [A.5.5](#)). This places our procedure at a chosen point on the frontier that trades the false-positive rate against the match rate, in the spirit of the comparison in [Abramitzky et al.](#)

[2021].

For each ambiguous pair we present a large language model (Claude Haiku 4.5) with a structured summary of the two records—the scientist’s name, field, affiliation history, and representative paper titles, and the donor’s name, state, employer, occupation, most recent active cycle, and CF score—and ask whether they refer to the same individual. The prompt instructs the model that scientists are geographically and institutionally mobile, so that a true match may disagree on location or employer, while a shared rare employer or field constitutes strong positive evidence; this mirrors the treatment of mobility elsewhere in our procedure. The model returns a structured verdict—a same-person judgment, a confidence in $[0, 1]$, and a one-sentence rationale—and we accept a band pair only when the judgment is affirmative *and* the reported confidence is at least 0.90. We calibrated this rule on the manually labeled audit before applying it.

We match a total of 125,284 scientists. The vast majority of these matches (116,860) are linked directly as the posterior probability is above the acceptance threshold. Using the LLM on marginal cases yields another 8,424 scientist matches.

A.5.5 Validation

We evaluate the procedure against a manually linked benchmark. We sampled scientists from the Dimensions file, oversampling successful and unique links. For each random pair, we conducted manual searches of the web and of DIME to determine the true donor match. Treating these as ground truth, we report the key match statistics in Table S1

Table S1: Summary Statistics of Matching Audit

	Linked	Not linked	
Truth: match	512	160	Recall 0.762
Truth: non-match	27	231	Specificity 0.895
	Precision 0.950	NPV 0.591	Accuracy 0.799

Notes: The table presents a confusion matrix summarizing the performance of our linking procedure against a manually verified benchmark. An observation is a scientist in the audit sample for whom we manually searched the web and DIME to establish the true donor match (if any). Rows classify scientists by ground truth (whether a true donor match exists); columns classify them by whether our procedure linked them to a donor. The diagonal cells (512 true matches linked, 231 true non-matches left unlinked) are correct classifications; the off-diagonal cells (27 false positives, 160 false negatives) are errors. Recall is the share of true matches that are linked; specificity is the share of true non-matches left unlinked; precision is the share of linked pairs that are true matches (the false-discovery rate is $1 - \text{Precision} = 0.050$); NPV is the share of unlinked scientists that are true non-matches; and accuracy is the overall share of correct classifications. The sample contains 930 scientists.

A.5.6 Discussion of potential biases from measurement of paper ideology

Noisy Measure: CF Score is an imperfect proxy for scientist ideology for two reasons. First, noise in the linking procedure creates some false positives. Second, our measure of paper slant only considers the slant of authors who have CF scores. The linked authors’ influence on paper ideology might be overwhelmed by the influence of the rest of the team. We expect such noise in our measure of scientist ideology to attenuate our estimates of alignment preferences towards zero.

Nonrepresentative sample: All scientists in our linked sample have donated to a political campaign. Given that the decision to donate to political campaigns is unlikely to be random, our sample may not be representative of the population of scientists. In particular, these scientists may have stronger ideological beliefs than the median scientist.

B Descriptive statistics

B.1 Media outlets


We begin by providing concrete examples of outlets in our sample. Table S2 lists the 10 outlets that covered the most articles. The largest outlets in our sample include large news aggregators like Yahoo! News, national newspapers like New York Times, national magazines like Forbes, and international outlets like BBC. The majority of the outlets in Table S2 are left-leaning. Figure S1 shows that this is also true in the full sample.

Table S2: Examples of outlets in sample

Outlet	Scientific Articles Covered	BMA Slant
msn.com	118926	-0.08
news.yahoo.com	113376	0.05
finance.yahoo.com	26140	0.08
forbes.com	23250	0.06
dailymail.co.uk	23199	0.29
nytimes.com	20345	-0.55
theguardian.com	17965	-0.57
washingtonpost.com	17035	-0.26
bbc.co.uk	14322	-0.33
edition.cnn.com	12518	-0.26

Notes: The table presents the top ten outlets by the number of scientific articles covered. The first column lists the outlet name. The second column lists the number of scientific articles covered. The third column lists the outlet slant as measured by the BMA score.

Figure S1: Distribution of outlet ideology




Notes: The figure presents the histogram of outlets by outlet slant. The x -axis represents the outlet slant as measured by the BMA score. The y -axis represents the number of outlets.

B.2 Ideology of scientists and fields of science

In Figure S2, we present the distribution of the CF Scores of the linked scientists. Most scientists in our sample are liberal. The ideological distribution of scientists in our sample is similar to the ideological distribution of all scientists in the US [Bonica, 2014].


Figure S3 presents the mean ideology of papers across different scientific fields. Note that all fields, on average, are liberal-leaning. However, there is still considerable heterogeneity across fields. Fields like Business Science and Agriculture Science are more conservative relative to fields like Communications and Culture.

Figure S2: Distribution of scientist ideology



Notes: The figure presents the distribution of ideology for linked scientists using our donations-based measure. The x -axis represents the scientist's ideology as measured by the CF Score. The y -axis represents the number of scientists. The blue and red vertical lines present the CF Score of the median Democratic and the median Republic senator, respectively.

Figure S3: Mean Ideology of Scientists by Field



Notes: The figure presents differences in mean ideology across fields of science. The x -axis represents ideology as measured by the CF Score. The y -axis represents different scientific fields. The points denoted the mean CF score and the lines denote the 95% confidence intervals for each field.

C Linear probability models

We estimate models of the following type on the set of all **outlet-paper** pairs (o, p) :

$$Y_{op} = \beta a(I_o, I_p) + \gamma X_{op} + \epsilon_{op}, \quad (1)$$

where Y_{op} is an indicator for whether the paper was covered by the outlet at least once between 2015-2023, $a(I_o, I_p)$ is the ideological alignment between outlet ideology I_o and paper ideology I_p , X_{op} represents controls, and ϵ_{op} represents idiosyncratic errors. Our primary measure of alignment $a(I_o, I_p)$ is the absolute value of the difference in ideology scores for the paper and outlet $a(I_o, I_p) = |I_o - I_p|$. To aid with interpretation, we standardize the measure to have a mean of 0 and a standard deviation of 1. Our results are robust to alternative alignment measures (see Section D.1). The parameter β captures the association between alignment and coverage probability, conditional on controls X_{op} . In all specifications, our controls X_{op} include fixed effects that flexibly control for outlet-specific or field-specific features over time. The key difference between this analysis and the nonparametric analysis presented in the main text is that this analysis only allows for a linear effect of ideological alignment, while Figures 2 and 3 in the main text allow for a flexible relationship between outlet ideology, paper ideology, and coverage probability.

We present estimates in Table S3. Column (1) includes Field-Year and Outlet fixed effects. Columns (2) builds on the previous specification and includes more granular Field-Outlet fixed effect. In Columns (3) and (4), we control for non-political paper features using paper fixed effects and our quality index, respectively.

Table S3: Parametric estimates of alignment association (LPMs)

Dependent Variable: Model:	(1)	Coverage Probability		
		(2)	(3)	(4)
<i>Variables</i>				
Alignment	0.0009*** (4.33×10^{-5})	0.0009*** (4.32×10^{-5})	0.0005*** (5.91×10^{-5})	-5.66×10^{-5} (4.41×10^{-5})
Quality (PCA)				0.0065*** (2.97×10^{-5})
<i>Fixed-effects</i>				
Field-Year	Yes	Yes	Yes	Yes
Outlet	Yes		Yes	
Field-Outlet		Yes		Yes
Paper			Yes	
<i>Fit statistics</i>				
Observations	14,897,653	14,897,653	14,897,653	14,451,758
R ²	0.09140	0.09763	0.13296	0.10050
Dependent variable mean	0.01887	0.01887	0.01887	0.01896
<i>IID standard-errors in parentheses</i>				
<i>Signif. Codes: ***: 0.01, **: 0.05, *: 0.1</i>				

Notes: The table presents the parameter estimates of the linear probability model defined in Equation 1 for different sets of controls. An observation is an outlet-paper pair. The parameter of interest is β , which captures the extent to which ideological alignment a_{op} between an outlet and a paper is associated with coverage probability. Alignment $a(I_o, I_p)$ is the standardized absolute value of the difference in the ideology of the paper and outlet. We only include papers that were cited by at least one outlet.

D Robustness checks

D.1 Alternate alignment measures

In our main analysis, we define alignment between an outlet and a paper as the absolute value of the difference in the ideology $a(I_o, I_p) = |I_o - I_p|$. A potential measurement issue is that the BMA score and CF Scores are on different scales. Even though 0 represents the ideological center for both measures (and so the signs of the measures are consistent), the magnitudes are incommensurate. To assuage these concerns, we employ two other measures of alignment.

Our first alternate measure compares relative ideological positions. For each outlet, we estimate the quantile of their ideology in the distribution outlet ideologies $q(I_o)$. We apply the same procedure for scientific papers $q(I_p)$.³ Our relative alignment measure is the absolute value of the difference between these quantiles $a_q(I_o, I_p) = |q(I_o) - q(I_p)|$. By construction, the scales of the quantile measures are commensurate. In this measure, an outlet is perfectly aligned with a paper if both fall at the same point in their corresponding ideological distributions. Note that the ideological distribution of scientific ideology is left-skewed compared to that of outlet ideology. Therefore, a paper that is as conservative as an outlet on the relative scale will be more liberal on the absolute scale. We emphasize that there is no simple mapping between these two alignment measures as they capture different notions of alignment. Our second measure is a binary measure of alignment $a_b(I_o, I_p) = 1\{\text{sign}(I_o) = \text{sign}(I_p)\}$: an outlet and paper are aligned if they are both

³Formally, $\mathbb{P}(I > I_o) = q(I_o)$, where I is distributed according to the ideology distribution of outlets.

liberal (conservative); otherwise, they are not aligned. This measure only relies on the sign of the ideology and not the magnitude. This measure is invariant to any scale changes in the ideology measure. However, there is a loss of information with this measure, since it is a coarser measure of alignment.

Table [S4](#) presents estimates of alignment preferences using these alternative measures. The regression specification is that defined in Equation [1](#). We include field-year and field-outlet fixed effects in both specifications. We consistently find positive and quantitatively meaningful alignment preferences. These results are significant at all conventional levels. In summary, the association of alignment with coverage is not driven by any particular definition of alignment. As one would expect, the interpretation and magnitude of these alignment preferences do vary across these different definitions.

Table S4: Alignment associations using alternate alignment Measures

Dependent Variable: Model:	Coverage Probability	
	(1)	(2)
<i>Variables</i>		
Alignment Quantile (Relative Scale)	0.0031*** (0.0002)	
Aligned (Binary Scale)		0.0008*** (5.04×10^{-5})
<i>Fixed-effects</i>		
Field-Year	Yes	Yes
Field-Outlet	Yes	Yes
<i>Fit statistics</i>		
Observations	14,897,653	14,897,653
R ²	0.09763	0.09762
Dependent variable mean	0.01887	0.01887
<i>IID standard-errors in parentheses</i>		
<i>Signif. Codes: ***: 0.01, **: 0.05, *: 0.1</i>		

Notes: The table presents the parameter estimates of the linear probability model defined in Equation 1 for different sets of controls and measures of alignment. Each observation in the model is an outlet-paper pair. The parameter of interest is β which captures the extent to which ideological alignment a_{op} between an outlet and a paper is associated with coverage probability. In column (1), alignment $a_q(I_o, I_p)$ is the absolute value of the difference in the ideology quantile of the paper and outlet. In column (2), alignment $a_b(I_o, I_p)$ is an indicator which is equal to 1 if both the paper and outlet are liberal (conservative). We only include papers that were cited by at least one outlet. All columns include field-year and paper fixed effects.

D.2 Parametric quality analysis

We examine how coverage probability depends on alignment and different quality measures. We estimate the following linear LPMs

$$Y_{op} = \beta a(I_o, I_p) + \alpha q_p + \gamma X_{op} + \epsilon_{op}, \quad (2)$$

where Y_{op} is an indicator for whether the paper was covered by the outlet at least once between 2015-2023, $a(I_o, I_p)$ is the ideological alignment between outlet ideology I_o and paper ideology I_p , q_p is the quality of paper p , X_{op} represents controls, and ϵ_{op} represents idiosyncratic errors. We include field-year fixed effects and field-outlet fixed effects as controls. We examine four measures of quality: number of academic citations, amount of research grants, number of patent citations, and journal impact.

The results are presented in Table S5. For each individual quality measure, we find that coverage probability increases in quality. The association between quality measures and coverage probability is stronger than that of alignment and quality measures. This result reaffirms that quality affects coverage probability much more than alignment.

Table S5: Parametric Estimates of Alignment and Quality Association

Dependent Variable: Model:	(1)	(2)	Coverage Probability		
			(3)	(4)	(5)
<i>Variables</i>					
Alignment	0.0002*** (4.32×10^{-5})	0.0008*** (4.33×10^{-5})	0.0008*** (4.32×10^{-5})	0.0005*** (4.4×10^{-5})	-5.66×10^{-5} (4.41×10^{-5})
Academic Citations (log)	0.0082*** (3.55×10^{-5})				
Research Grants (log)		0.0001*** (4.27×10^{-6})			
Patents (log)			0.0045*** (7.01×10^{-5})		
Journal Impact (log)				0.0044*** (3.47×10^{-5})	
Quality (PCA)					0.0065*** (2.97×10^{-5})
<i>Fixed-effects</i>					
Field-Year	Yes	Yes	Yes	Yes	Yes
Field-Outlet	Yes	Yes	Yes	Yes	Yes
<i>Fit statistics</i>					
Observations	14,895,938	14,895,938	14,895,938	14,466,278	14,451,758
R ²	0.10084	0.09764	0.09785	0.09842	0.10050
Dependent variable mean	0.01886	0.01886	0.01886	0.01896	0.01896
<i>IID standard-errors in parentheses</i>					
<i>Signif. Codes: ***: 0.01, **: 0.05, *: 0.1</i>					


Notes: The table presents the parameter estimates of the linear probability model defined in Equation 2 for different quality measures. Each observation in the model is an outlet-paper pair. Alignment $a_q(I_o, I_p)$ is the standardized absolute value of the difference in the ideology of the paper and outlet. The quality measures are number of academic citations, amount of research grants, number of patent citations, and journal impact (h -index) for columns (1), (2), (3), and (4), respectively. We only include papers that were cited by at least one outlet. All columns include field-year and field-outlet fixed effects.

D.3 Quality association by field of science

In Figure S4 we present the association between coverage propensity and alignment alongside the association between coverage propensity and the quality index across fields of science. For each field, these associations are estimated simultane-

ously using an LPM as in Equation 2. For all fields, the coefficient on the quality index is larger than the coefficient on alignment.

Figure S4: Heterogeneity in alignment and quality associations by field of science



Notes: The figure presents the relationship between alignment and quality with media coverage for different fields of science. The y -axis represents a field of science. The x -axis presents the association between coverage propensity and alignment as well as quality, denoted by different colors. The alignment measure and quality index are normalized to have a mean of 0 and a standard deviation of 1.

References

- Ran Abramitzky, Leah Boustan, Katherine Eriksson, James Feigenbaum, and Santiago Pérez. Automated linking of historical data. *Journal of Economic Literature*, 59(3):865–918, September 2021. doi: 10.1257/jel.20201599.
- Eytan Bakshy, Solomon Messing, and Lada A Adamic. Exposure to ideologically diverse news and opinion on facebook. *Science*, 348(6239):1130–1132, 2015.
- Adam Bonica. Mapping the ideological marketplace. *Am. J. Pol. Sci.*, 58(2):367–386, 2014.
- Adam Bonica. Database on ideology, money in politics, and elections. *Public version 4.0 [Computer file]*. Stanford, CA: Stanford University Libraries., 2024. URL <https://data.stanford.edu/dime>.
- Ted Enamorado, Benjamin Fifield, and Kosuke Imai. Using a probabilistic model to assist merging of large-scale administrative records. *American Political Science Review*, 113(2):353–371, 2019. doi: 10.1017/S0003055418000783.
- Ivan P. Fellegi and Alan B. Sunter. A theory for record linkage. *Journal of the American Statistical Association*, 64(328):1183–1210, 1969. doi: 10.1080/01621459.1969.10501049.
- Christian Herzog, Daniel Hook, and Stacy Konkiel. Dimensions: Bringing down barriers between scientometricians and data. *Quantitative Science Studies*, 1(1): 387–395, 2020.